

A Lean tactic for normalising ring expressions with exponents

Anne Baanen

Lean Forward
Vrije Universiteit Amsterdam

IJCAR, 2 July 2020



Lean is a proof assistant based on the calculus of constructions [2]. It has a simple kernel for proof checking and an elaborator with powerful tactic support.

The Lean community is developing `mathlib`, a repository of formalised classical mathematics proofs and proof automation.

The Lean Forward project aims to make proof assistants accessible to mathematicians by developing proof automation informed by users' needs.

To a mathematician, the following is obvious:

$$2^{n+1} - 1 = 2 * 2^n - 1$$

Lean should do it automatically.

The previous `ring` tactic [1] uses Horner normal form: efficient for the semiring operators `+` and `*`, but it doesn't support exponentiation `^`.

To a mathematician, the following is obvious:

$$2^{n+1} - 1 = 2 * 2^n - 1$$

Lean should do it automatically.

The previous `ring` tactic [1] uses Horner normal form: efficient for the semiring operators `+` and `*`, but it doesn't support exponentiation `^`.

It's not too hard to solve this manually: rewrite $a^{n+1} = a^n * a$ and `ring` can finish by applying commutativity.

Such rules don't work unconditionally: x^{100} should not become 100 multiplications of `x`.

Goal: a practical normalising tactic for expressions with $+$, $*$ and $^$, numerals (in \mathbb{Q}) and variables. It should solve all goals that `ring` can and be approximately as fast.

Goal: a practical normalising tactic for expressions with $+$, $*$ and $^{\wedge}$, numerals (in \mathbb{Q}) and variables. It should solve all goals that `ring` can and be approximately as fast.

To prove $a = b$, normalise a giving $p_a : a = a'$ and normalise b giving $p_b : b = b'$, then check a' is identical to b' . If identical, $p_a p_b^{-1}$ proves $a = b$.

Lean tactics typically construct proof terms directly (no reflection), since the elaborator is faster than the kernel.

The normal form is a syntax tree in the type family `ex`. The children for each node are restricted by a parameter `ex_type`:

```
inductive ex_type : Type
| sum | prod | exp | base
inductive ex : ex_type → Type
| zero   : ex_info → ex sum -- 0
| sum    : ex_info → ex prod → ex sum → ex sum -- +
| coeff  : ex_info → coeff → ex prod -- rat
| prod   : ex_info → ex exp → ex prod → ex prod -- *
| exp    : ex_info → ex base → ex prod → ex exp -- ^
| var    : ex_info → atom → ex base -- atom
| sum_b  : ex_info → ex sum → ex base -- (...)
```

- $(a + b) + c$ is not allowed: left argument to `sum` must be a product
- $a * (b + c)$ is not allowed: right argument to `prod` must be a product

Commutativity: pick a linear order \prec on ex. Then sort $a + b + c + \dots$ so that $a \prec b \prec c \prec \dots$.

Commutativity: pick a linear order \prec on ex. Then sort $a + b + c + \dots$ so that $a \prec b \prec c \prec \dots$.

To limit expression size, don't unfold $100 * a$ to $a + a + \dots + a$.
This means keeping track of coefficients.

The function `add_overlap` decides when to add coefficients:

$$\begin{aligned}\text{add_overlap}(3 * x^2)(7 * x^2) &= 10 * x^2 \\ \text{add_overlap}(3 * x^2)(7 * y^2) &= 3 * x^2 + 7 * y^2 \\ \text{add_overlap}(3 * x^2)(-3 * x^2) &= 0 && (\text{not } 0 * x^2)\end{aligned}$$

Commutativity: pick a linear order \prec on ex. Then sort $a + b + c + \dots$ so that $a \prec b \prec c \prec \dots$.

To limit expression size, don't unfold $100 * a$ to $a + a + \dots + a$.
This means keeping track of coefficients.

The function `add_overlap` decides when to add coefficients:

$$\begin{aligned}\text{add_overlap } (3 * x^2) (7 * x^2) &= 10 * x^2 \\ \text{add_overlap } (3 * x^2) (7 * y^2) &= 3 * x^2 + 7 * y^2 \\ \text{add_overlap } (3 * x^2) (-3 * x^2) &= 0 \quad (\text{not } 0 * x^2)\end{aligned}$$

In a general semiring R , exponentiation has type $^{\wedge} : R \rightarrow \mathbb{N} \rightarrow R$.
During execution, `ring_exp` keeps track of the current type using a reader monad transformer.

To be practical, the `ring_exp` tactic must be fast.

The Horner form used by the `ring` tactic is optimal for $+$ and $*$.

Optimisation is needed to achieve acceptable running time.

To be practical, the `ring_exp` tactic must be fast.

The Horner form used by the `ring` tactic is optimal for $+$ and $*$.

Optimisation is needed to achieve acceptable running time.

Typeclass instances and implicit arguments cost time to infer,
so they are cached as much as possible:

instances are stored with the current type in the reader monad,
implicit arguments and intermediate values in the `ex_info` field of `ex`.

Optimisations

To be practical, the `ring_exp` tactic must be fast.

The Horner form used by the `ring` tactic is optimal for $+$ and $*$.

Optimisation is needed to achieve acceptable running time.

Typeclass instances and implicit arguments cost time to infer,
so they are cached as much as possible:

instances are stored with the current type in the reader monad,
implicit arguments and intermediate values in the `ex_info` field of `ex`.

Result: on average a constant factor slowdown compared to `ring`.

Noticeably faster on problems with larger exponents

(20 times on $x^{50} * x^{50} = x^{100}$), also in practice for

$$(1 + x^2 + x^4 + x^6) * (1 + x) = 1 + x + x^2 + x^3 + x^4 + x^5 + x^6 + x^7.$$



Benjamin Grégoire and Assia Mahboubi. “Proving equalities in a commutative ring done right in Coq”. In: TPHOLs 2005. Ed. by J. Hurd and T. Melham. Vol. 3603. LNCS. Springer, Berlin, Heidelberg, 2005, pp. 98–113. DOI: [10.1007/11541868_7](https://doi.org/10.1007/11541868_7).



Leonardo de Moura et al. “The Lean theorem prover (system description)”. In: Automated Deduction - CADE-25. Ed. by A. P. Felty and A. Middeldorp. Vol. 9195. LNCS. Springer, Cham, 2015, pp. 378–388. DOI: [10.1007/978-3-319-21401-6_26](https://doi.org/10.1007/978-3-319-21401-6_26).